

# Linear Regression

DR. SETHAVIDH GERTPHOL

# Today's Outline

- **Linear Regression**
- **Polynomial Features**
- **Gradient Descent**

# Linear Regression

- Regression: predict the numerical output using input features
- Linear Regression: use this equation for prediction

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

- Note that there are  $n$  features, but  $n+1$  model parameters
  - $\theta_0$  is bias,  $\theta_i$  is coefficient of  $x_i$
- When there is only one feature, the equation has the same form as  $y = b + ax$

- $\hat{y}$  is the predicted value.
- $n$  is the number of features.
- $x_i$  is the  $i^{\text{th}}$  feature value.
- $\theta_j$  is the  $j^{\text{th}}$  model parameter,

# Finding the best thetas (model parameters)

- Which set of thetas is best?
- We must define a **cost function** and the best thetas are the ones that minimizes the cost
- Mean Square Error is a good cost function
  - Correlates with a metric used to evaluation regression model
  - Easy to optimize

$$\text{MSE}(\mathbf{X}, h_{\theta}) = \frac{1}{m} \sum_{i=1}^m \left( \theta^T \mathbf{x}^{(i)} - y^{(i)} \right)^2$$

Diagram illustrating the Mean Square Error (MSE) formula with annotations:

- Model parameters**: Points to  $\theta$  in the formula.
- Actual value of that sample**: Points to  $y^{(i)}$  in the formula.
- Prediction of one sample**: Points to  $\theta^T \mathbf{x}^{(i)}$  in the formula.
- m is the total number of samples**: Points to  $m$  in the denominator.

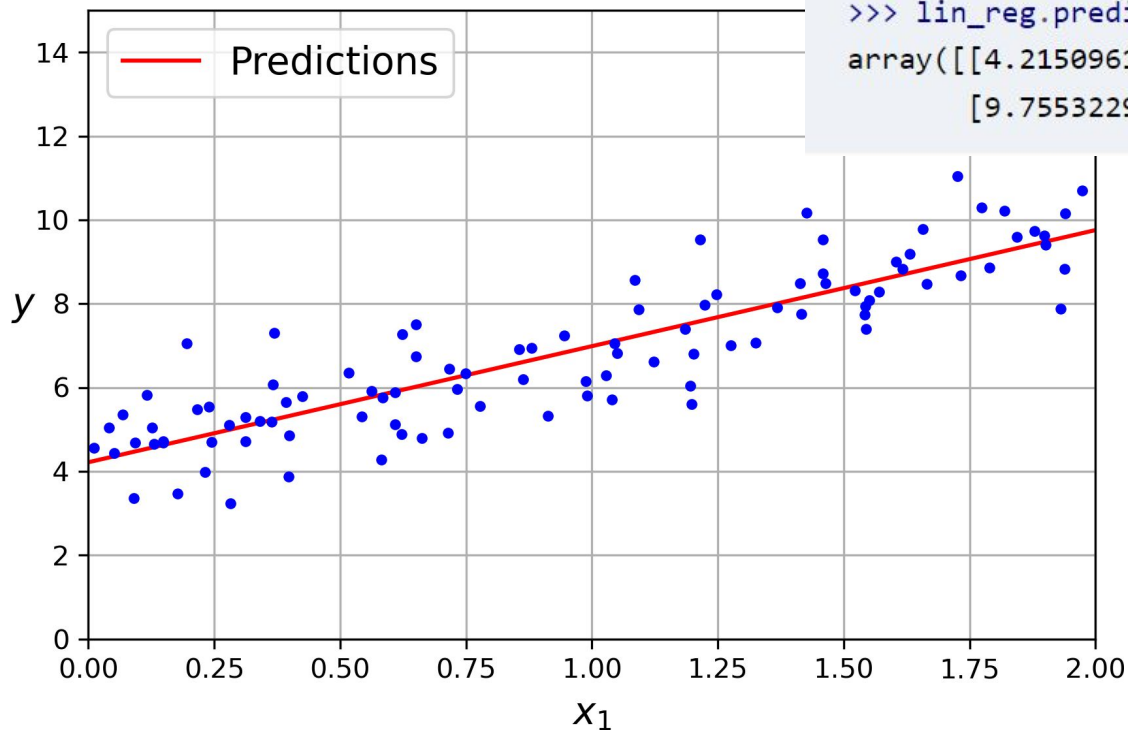
# Solving for thetas

- There is a closed-form solution for MSE cost function called the **Normal Equation**

$$\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- Sklearn uses **SVD (Singular Value Decomposition)** to calculate the pseudoinverse matrix
- This is a calculation, so there is **no hyperparameter** to tune
- No need for feature scaling
- Use one-hot encoding to transform categorical features

# In Sklearn



import

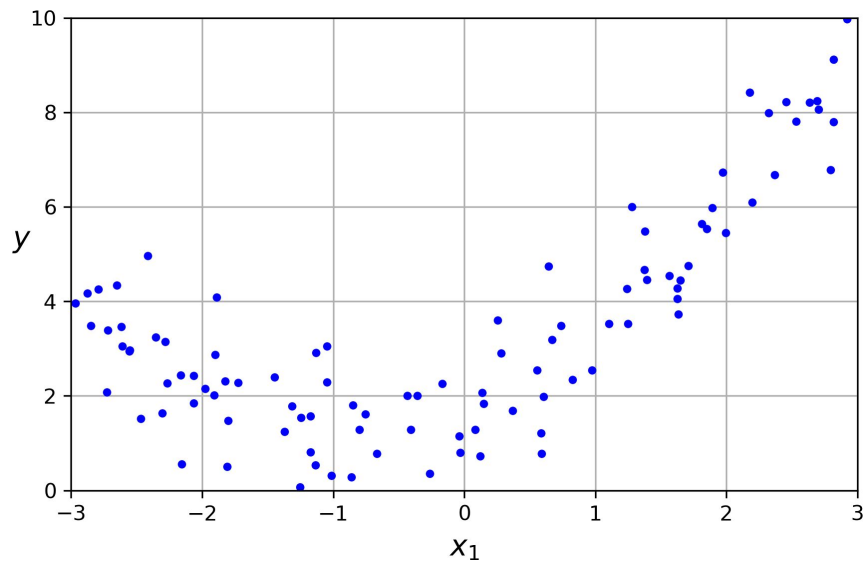
```
>>> from sklearn.linear_model import LinearRegression
>>> lin_reg = LinearRegression()
>>> lin_reg.fit(X, y)
>>> lin_reg.intercept_, lin_reg.coef_
(array([4.21509616]), array([[2.77011339]]))
>>> lin_reg.predict(X_new)
array([[4.21509616],
       [9.75532293]])
```

$\theta_0$  is in **intercept\_**,  
other  $\theta_i$  are in **coef\_**

Note that `LinearRegression()`  
adds the bias **automatically**

# Polynomial Regression

- What if the data has underlying polynomial terms?
- Can add polynomial terms to feature matrix and use LinearRegression to fit the new set of features



# PolynomialFeatures

import

```
>>> from sklearn.preprocessing import PolynomialFeatures
>>> poly_features = PolynomialFeatures(degree=2, include_bias=False)
>>> X_poly = poly_features.fit_transform(X)
>>> X[0]
array([-0.75275929])
>>> X_poly[0]
array([-0.75275929,  0.56664654])
```

Transform X

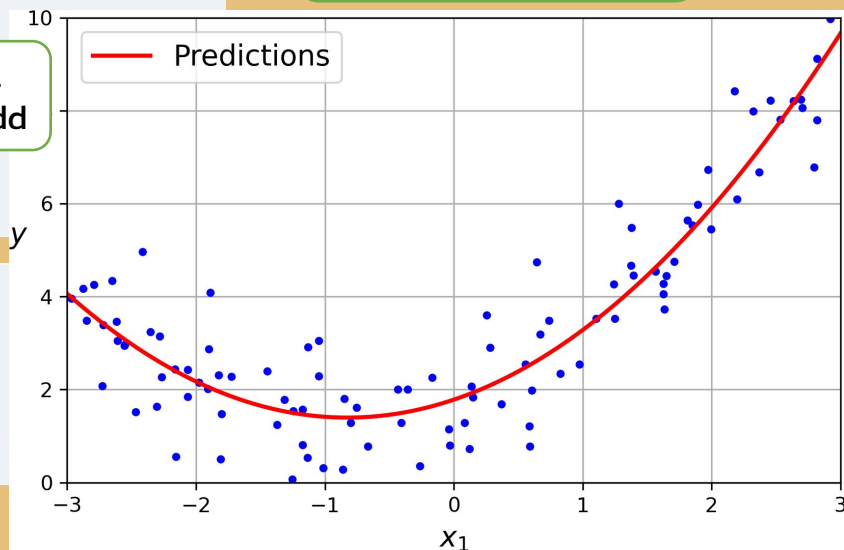
Polynomial  
degrees to add

No need to add bias  
because it will be added  
by LinearRegression

```
>>> lin_reg = LinearRegression()
>>> lin_reg.fit(X_poly, y)
>>> lin_reg.intercept_, lin_reg.coef_
(array([1.78134581]), array([[0.93366893, 0.56456263]]))
```

Important: PolynomialFeatures will  
**add all terms from lower degrees**  
and also **combination of terms** too

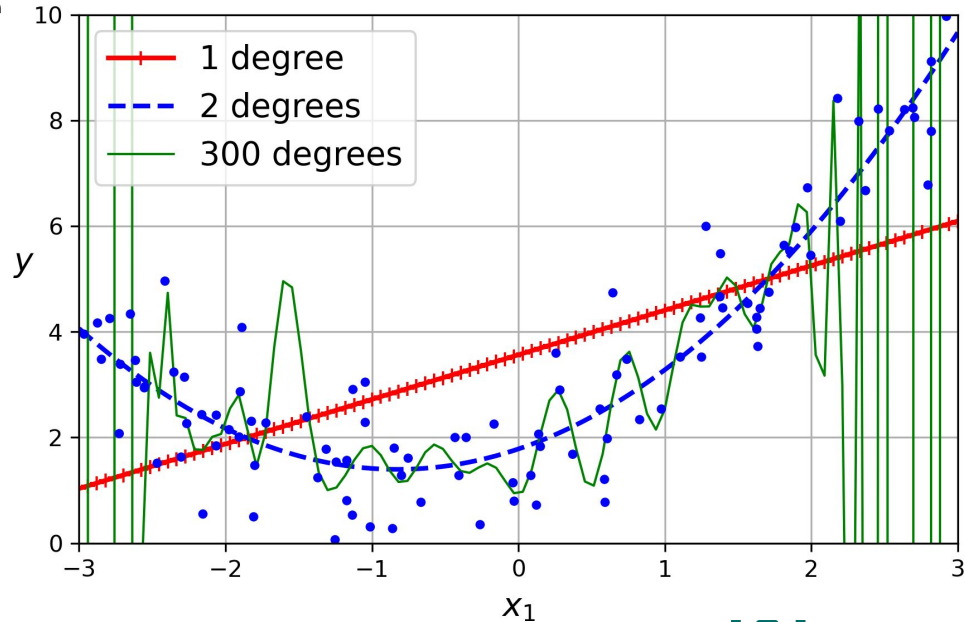
x, y with degree=3 will add:  $x^3, y^3, x^2, y^2, x^3, xy, xy^2, x^2y$





# What value of degree is best?

- Low degree: model is not flexible enough, may have high **bias**
  - Bias is a wrong assumption, such as data is linear but in fact it is polynomial
  - Usually leads to **underfitting**
- High degree: model is too complex, have high **variance**
  - Variance is sensitivity to small variations in training data
  - Usually leads to **overfitting**
- Bias-Variance Trade-off
- Use **cross-validation** to find best value

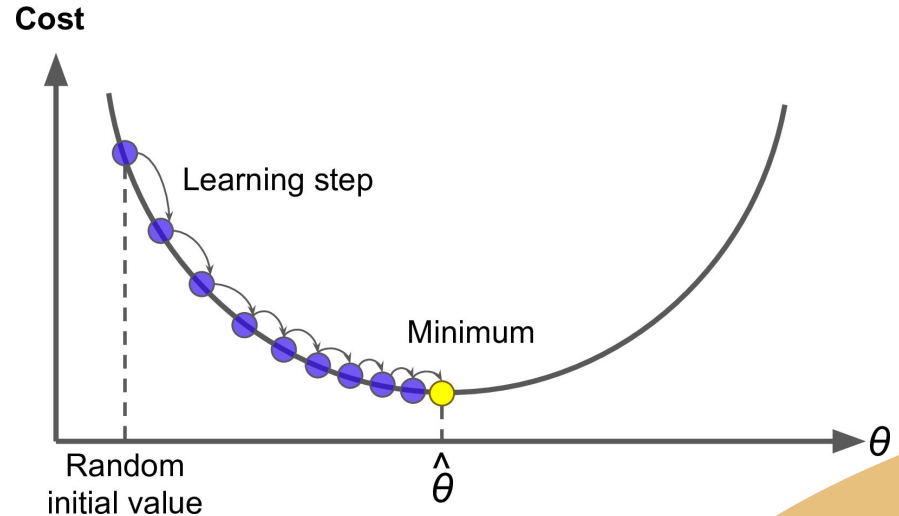


# Evaluation

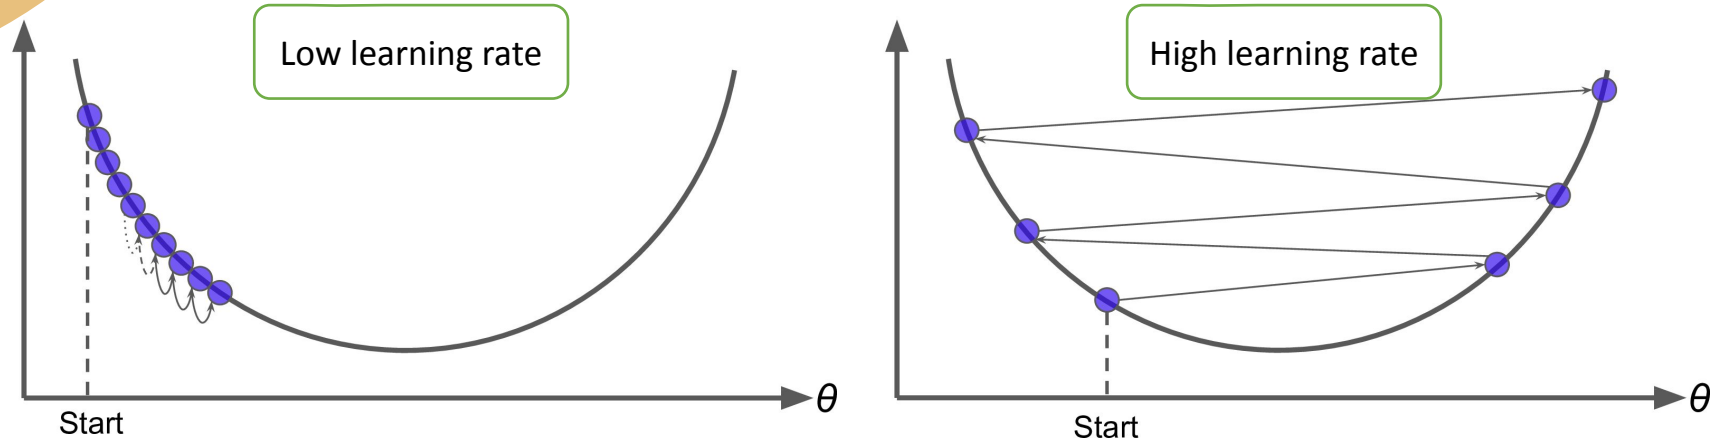
- Metrics that can be used (in sklearn.metrics)
  - mean\_absolute\_error
  - mean\_squared\_error
  - mean\_absolute\_percentage\_error
  - r2\_score

# Gradient Descent

- Normal Equation and SVD can take a long time with **large number of features**
- Gradient Descent
  - Optimization technique
  - **Start with random values** of thetas and **adjust them step-by-step** to reduce cost function
  - **Evaluating gradient** (slope) of cost function at current thetas and **follow the greatest descent** (steepest slope downward)
  - When **gradient is zero**, we have found thetas that give **minimum cost**



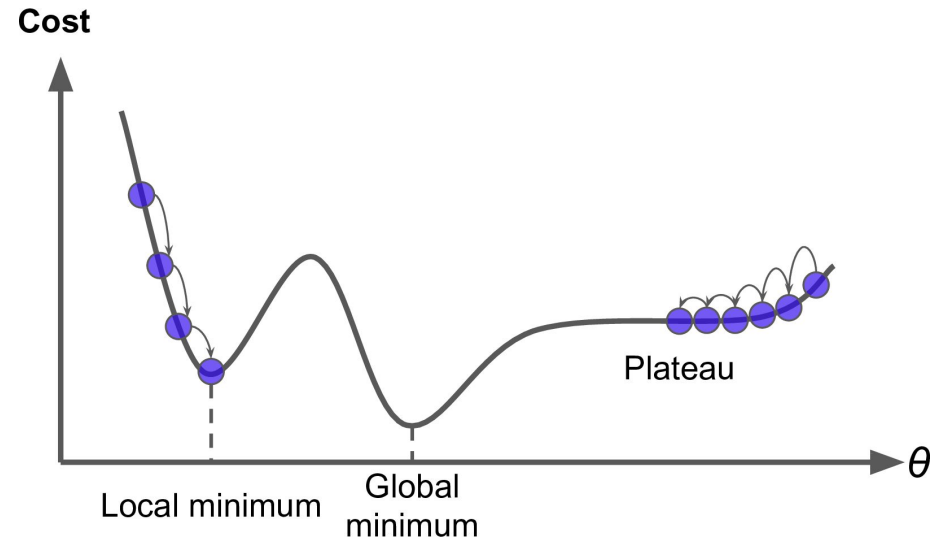
# Learning rate



- How much we adjust thetas during each step of gradient descent
- Trade-off:
  - Low learning rate: small step, take too long to find minimum
  - High learning rate: large step, may go over minimum and diverge

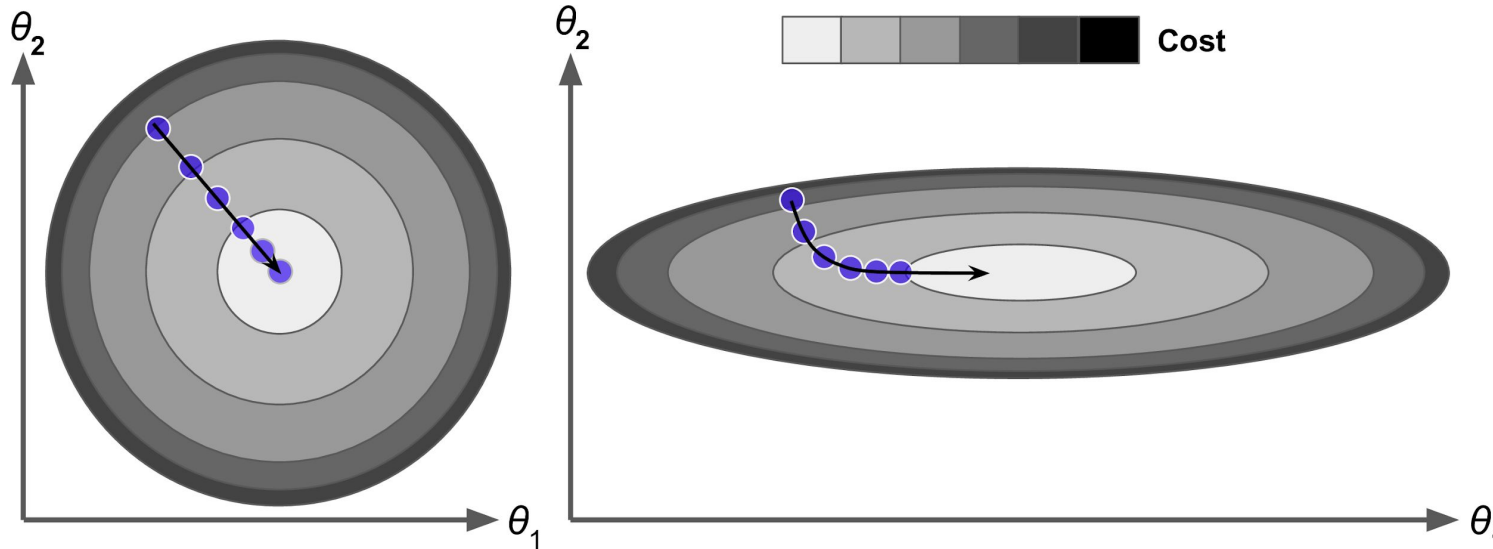
# Local and Global minima

- Not all cost functions have nice smooth graph with one minimum
- Many have several minima, ridges, peaks, plateau
- Gradient descent **may stop at local minimum** instead of global one
- May **not be able to escape local minimum** due to peak next to it
- May **take a long time to cross plateau**



# Feature scaling

- MSE cost function has a nice bowl shape with one minimum
- But if features are on different scale, may take long time to converge
- Should use **feature scaling** (e.g., StandardScaler) with Gradient Descent



# Batch Gradient Descent

- Calculate partial derivative of each theta on the cost function to find gradient
- Use gradient to adjust thetas
- Fast to calculate even with large number of features (thetas)
- Slow with large number of samples

$$\text{MSE}(\mathbf{X}, h_{\theta}) = \frac{1}{m} \sum_{i=1}^m \left( \theta^T \mathbf{x}^{(i)} - y^{(i)} \right)^2$$

$$\frac{\partial}{\partial \theta_j} \text{MSE}(\theta) = \frac{2}{m} \sum_{i=1}^m \left( \theta^T \mathbf{x}^{(i)} - y^{(i)} \right) x_j^{(i)}$$

$$\nabla_{\theta} \text{MSE}(\theta) = \begin{pmatrix} \frac{\partial}{\partial \theta_0} \text{MSE}(\theta) \\ \frac{\partial}{\partial \theta_1} \text{MSE}(\theta) \\ \vdots \\ \frac{\partial}{\partial \theta_n} \text{MSE}(\theta) \end{pmatrix} = \frac{2}{m} \mathbf{X}^T (\mathbf{X}\theta - \mathbf{y})$$

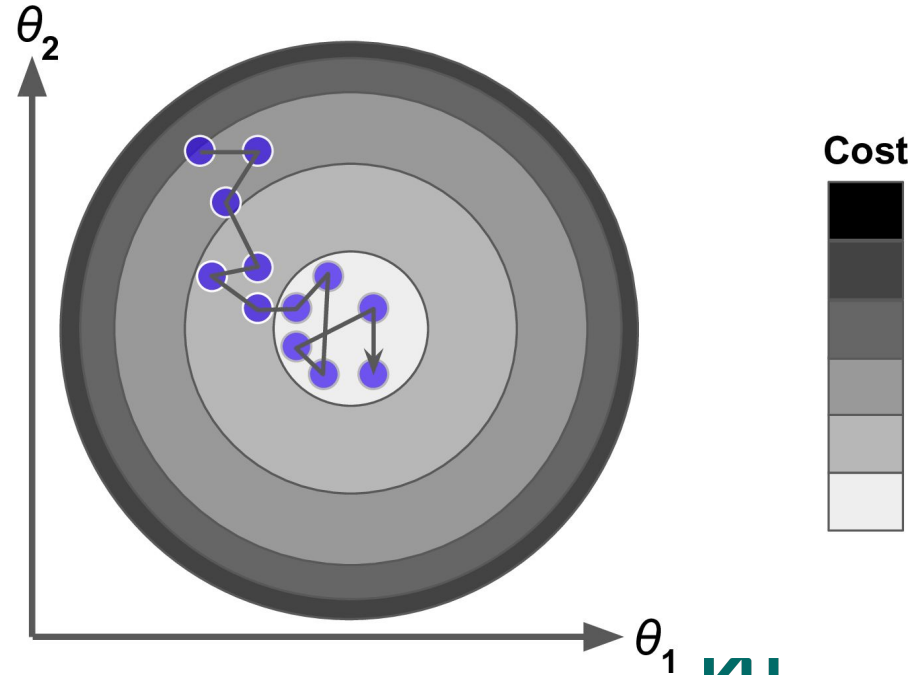
Cost function  
(MSE)

Partial derivative  
of  $\theta_j$

Matrix form of  
partial derivative  
of all thetas

# Stochastic Gradient Descent

- **Randomly pick 1 sample**, and calculate gradient based on that sample
- Very fast even with large number of samples
- Cost value may bounce around or even go up during one iteration
- Over time it will go near the minimum value, but never stops





# Stopping SGD (Stochastic Gradient Descent)

- Set the number of epochs to optimize
  - epoch: 1 round of  $m$  iterations where  $m$  is the number of samples
- Reduce learning rate in later epochs/iterations
  - Learning schedule: a function to adjust learning rate
- Change in cost:
  - If cost change is less than given threshold over later iterations

# SGD in Sklearn

```
from sklearn.linear_model import SGDRegressor

sgd_reg = SGDRegressor(max_iter=1000, tol=1e-5, penalty=None, eta0=0.01,
                        n_iter_no_change=100, random_state=42)

sgd_reg.fit(X, y.ravel()) # y.ravel() because fit() expects 1D targets
```

import

iterations

Regularization  
(not taught)

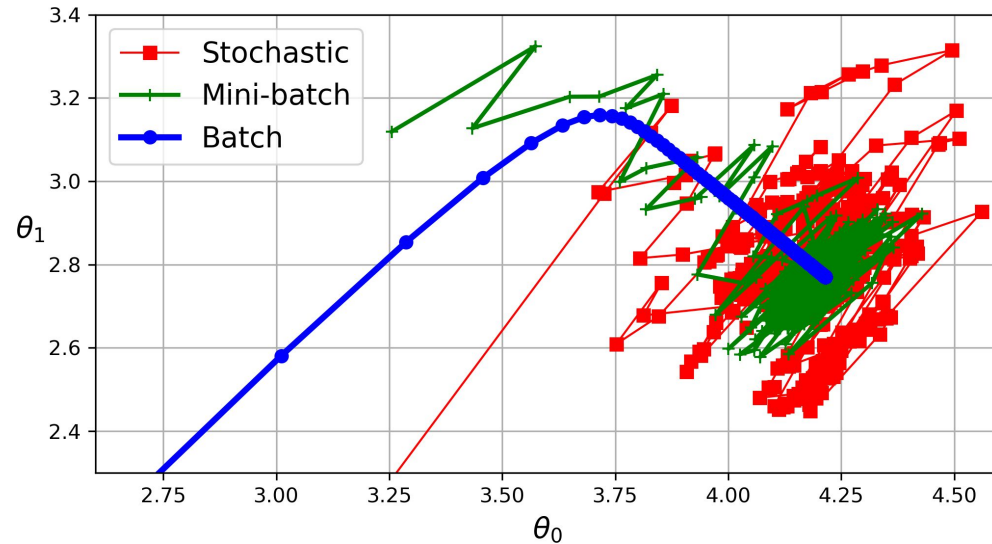
Starting  
learning rate  
(default  
schedule)

Stop when reach this  
number of iterations  
with less than tol change

```
>>> sgd_reg.intercept_, sgd_reg.coef_  
(array([4.21278812]), array([2.77270267]))
```

# Mini-Batch Gradient Descent

- Calculation of gradient
  - Batch: all samples
  - Stochastic: 1 random sample
  - Mini-batch: random subset of samples
- Progression of thetas in these 3 methods



# Things not taught

- **Learning Curve**: a technique to evaluate overfitting or underfitting using training and validation result (cross-validation) over iterations
- **Regularization**: techniques to reduce overfitting by introducing penalty term to cost function
- Regularized Linear Model
  - Ridge Regression
  - Lasso Regression
  - Elastic Net

# Reference

- Aurélien Géron, "Hands-On Machine Learning with Scikit-Learn and TensorFlow", O'Reilly Media, Inc., March 2017.