

Ensemble Method

DR. SETHAVIDH GERTPHOL

Today's Outline

- Ensemble Method
- Voting Classifier
- Bagging (Random Forest)
- Boosting
 - Adaboost
 - Gradient Boosting
- Stacking

Wisdom of the crowd

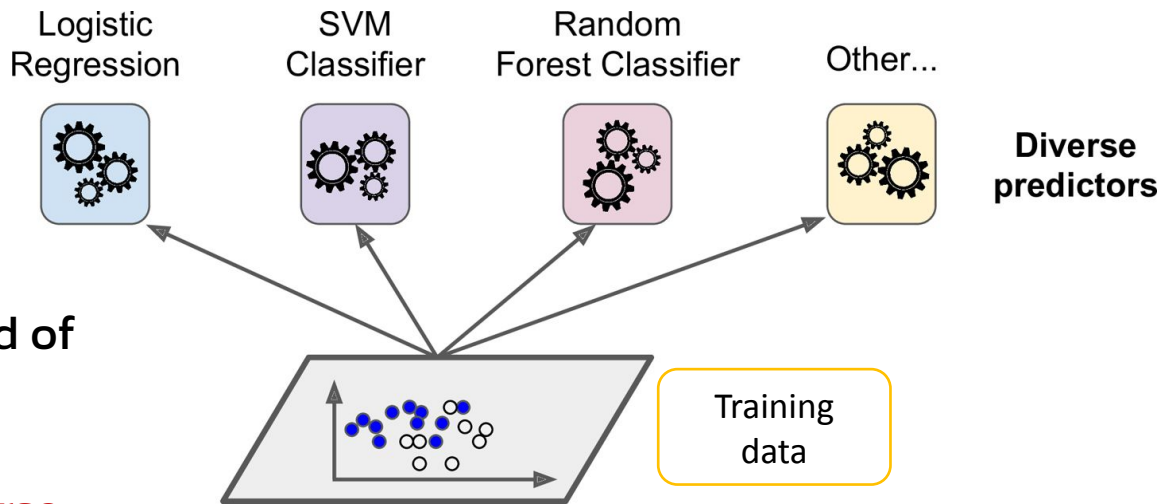
- Sir Francis Galton observed that an aggregated of estimates made by 787 persons about weight of an ox had **only 1% error** from the true wright.
- It may be better to **aggregate predictions from several models** instead of building a perfect model, even though some of them may be very bad.
- This is called **Ensemble Method**

Distribution of the estimates of the dressed weight of a particular living ox, made by 787 different persons.

Degrees of the length of Array 0°—100°	Estimates in lbs.	Centiles		Excess of Observed over Normal
		Observed deviates from 1207 lbs.	Normal p.e = 37	
5	1074	- 133	- 90	+ 43
10	1109	- 98	- 70	+ 28
15	1126	- 81	- 57	+ 24
20	1148	- 59	- 46	+ 13
q_1 25	1162	- 45	- 37	+ 8
30	1174	- 33	- 29	+ 4
35	1181	- 26	- 21	+ 5
40	1188	- 19	- 14	+ 5
45	1197	- 10	- 7	+ 3
m 50	1207	0	0	0
55	1214	+ 7	+ 7	0
60	1219	+ 12	+ 14	- 2
65	1225	+ 18	+ 21	- 3
70	1230	+ 23	+ 29	- 6
q_3 75	1236	+ 29	+ 37	- 8
80	1243	+ 36	+ 46	- 10
85	1254	+ 47	+ 57	- 10
90	1267	+ 52	+ 70	- 18
95	1293	+ 86	+ 90	- 4

q_1 , q_3 , the first and third quartiles, stand at 25° and 75° respectively.
 m , the median or middlemost value, stands at 50°.
 The dressed weight proved to be 1198 lbs.

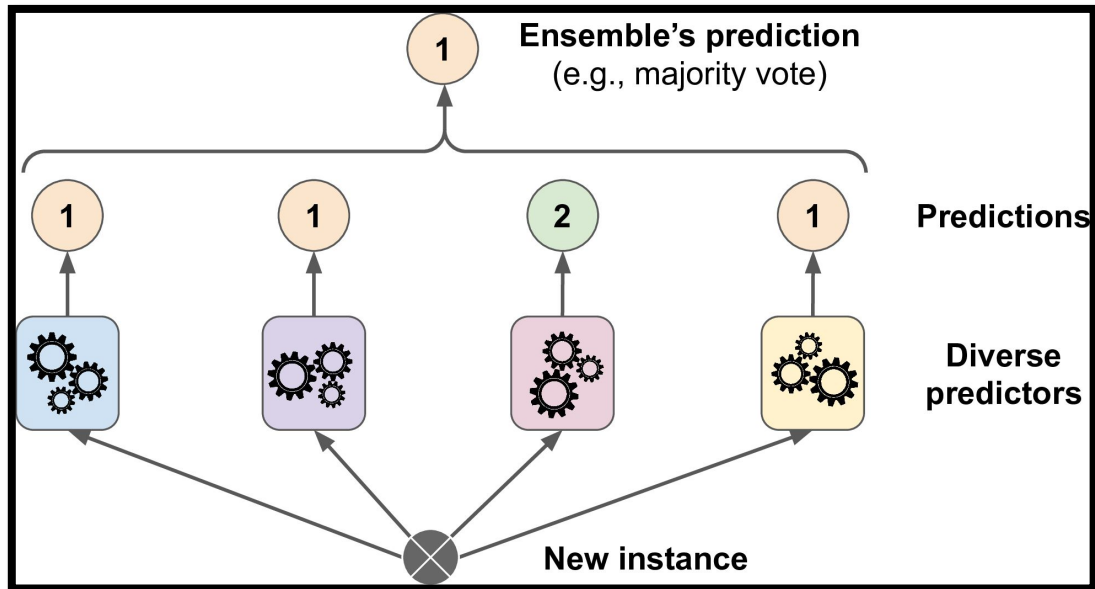
Ensemble Method



- Usually used at the end of project when we have several good models
- Models should be **diverse**
 - Using different algorithms or
 - Using different data

Voting Classifier

- A simple majority vote classifier
- Ensemble's prediction follows the majority vote of all classifiers



In Sklearn

```
from sklearn.datasets import make_moons
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC

X, y = make_moons(n_samples=500, noise=0.30, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)

voting_clf = VotingClassifier(
    estimators=[
        ('lr', LogisticRegression(random_state=42)),
        ('rf', RandomForestClassifier(random_state=42)),
        ('svc', SVC(random_state=42))
    ]
)
voting_clf.fit(X_train, y_train)
```

import

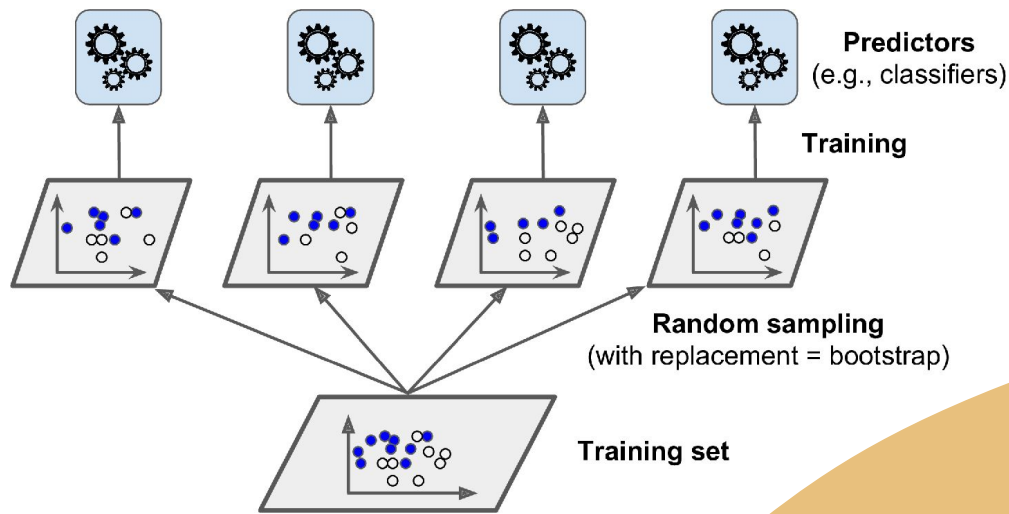
Generating
dataset

Each predictor is a
tuple of name and
predictor object

List of
predictors

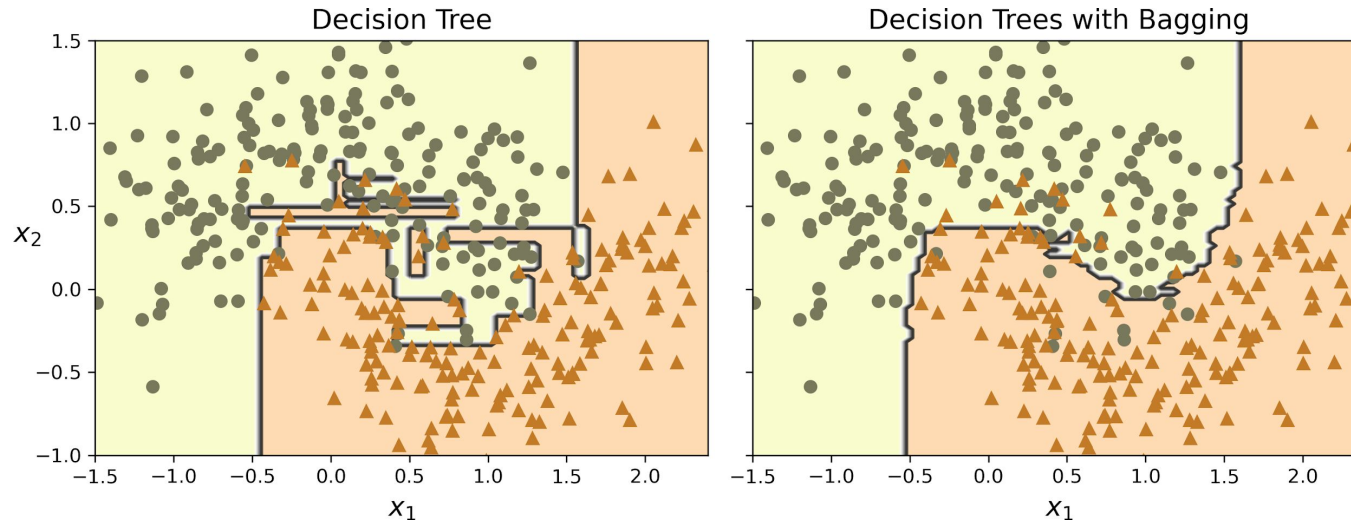
Bagging and Pasting

- One way to have diverse predictors is to use the **same algorithm** on **different training data**
- Use sampling to build different training set for each predictor
- **Bagging** (bootstrap aggregation) is **sampling with replacement**
- **Pasting** is sampling **without replacement**



Bagging

- Generally better performance and more preferred than pasting
- If sampling size of a predictor is equal to training size, one predictors on average will see only 67% of the training data (33% are duplicates)
- Creates **more bias** on each predictor, but this bias make predictors **more diverse**
- **Out-of-bag evaluation**: use samples that are not included in a sample set as a test set of that predictor (no need for separate validation set)



In Sklearn

import

```
from sklearn.ensemble import BaggingClassifier  
from sklearn.tree import DecisionTreeClassifier
```

Predictor ที่จะใช้ ใส่
option ที่ต้องการ

จำนวน predictor
ที่จะสร้าง

```
bag_clf = BaggingClassifier(DecisionTreeClassifier(), n_estimators=500,  
                           max_samples=100, random_state=42)  
bag_clf.fit(X_train, y_train)
```

จำนวน predictor
ที่จะสร้าง

```
>>> bag_clf = BaggingClassifier(DecisionTreeClassifier(), n_estimators=500,  
...                           oob_score=True, random_state=42)  
...  
>>> bag_clf.fit(X_train, y_train)  
>>> bag_clf.oob_score_  
0.896
```

Option เพื่อให้คำนวณ
out-of-bag score ด้วย

Random Forest

- An **ensemble of Decision Tree** built using bagging or pasting, usually with sample size equal to training set size
- Introduce additional randomness by selecting the best feature **among a random subset of features** to split each node
- Feature importance is the **weight average** of reduction in impurity of a node when using that feature to split it
 - Average across all trees
 - Weighted by the number of samples associated with that node

In Sklearn

import

```
from sklearn.ensemble import RandomForestClassifier

rnd_clf = RandomForestClassifier(n_estimators=500, max_leaf_nodes=16,
                                n_jobs=-1, random_state=42)

rnd_clf.fit(X_train, y_train)

y_pred_rf = rnd_clf.predict(X_test)
```

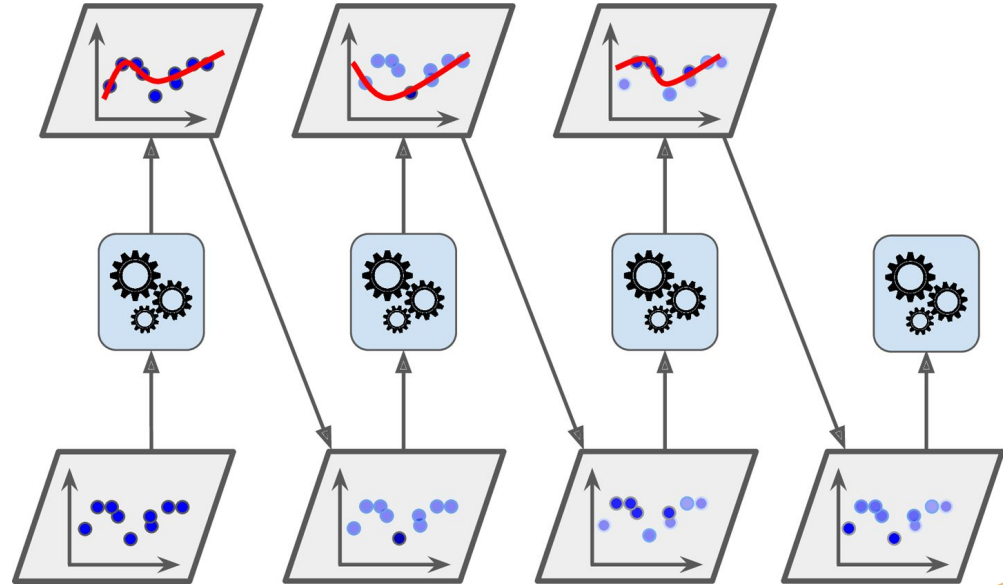
จำนวน predictor
ที่จะสร้าง

จำนวน CPU core ที่
ต้องการใช้
-1 คือใช้ทุก core

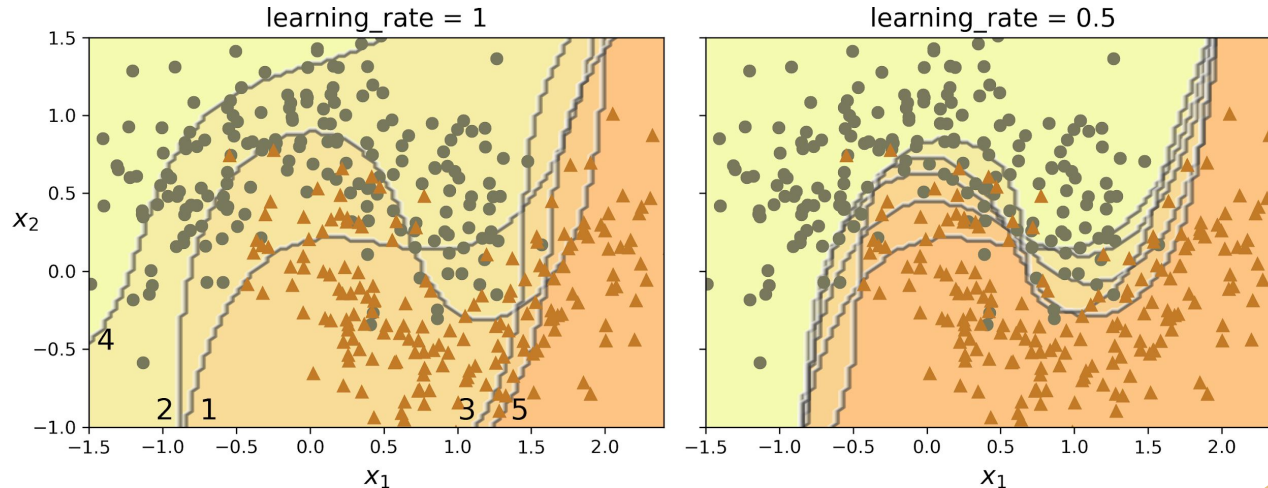
Option อื่นใช้เหมือนของ
DecisionTreeClassifier

Boosting

- An ensemble method that **trains predictors sequentially**, with later ones **trying to improve on the previous ones**
- AdaBoost
 - Train 1 predictor, see which **training instances** are predicted wrong
 - **Increase weight** of those instances, train another predictors
 - Keep doing until reaching the number of predictors required



- **Learning rate**: affect how much weight of wrong instances got boosted
- **Predictor's weight**: depends on how many instances it got right
- Final prediction is the **weighted average of all predictors' prediction**



In Sklearn

import

Predictor ที่จะใช้

Learning rate

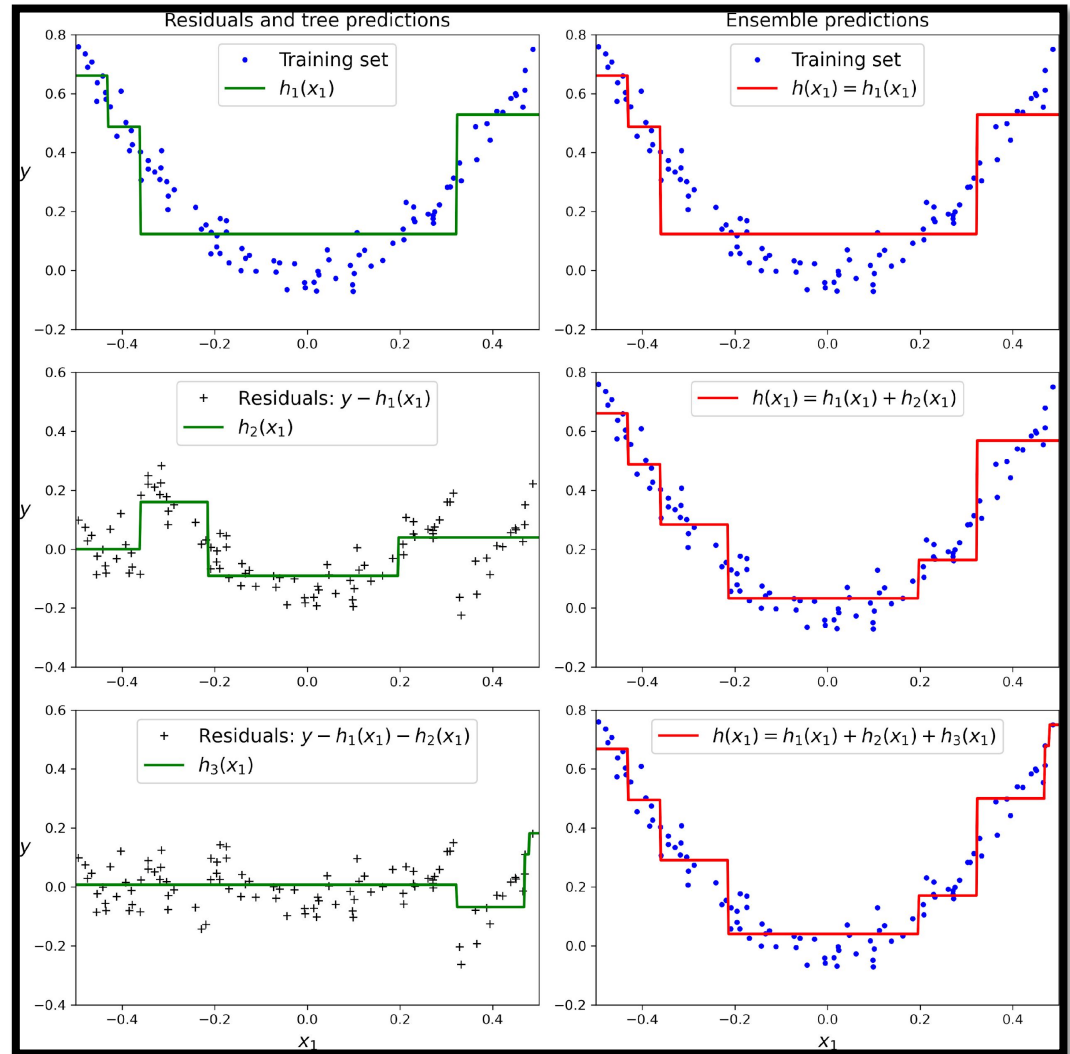
```
from sklearn.ensemble import AdaBoostClassifier
```

```
ada_clf = AdaBoostClassifier(  
    DecisionTreeClassifier(max_depth=1), n_estimators=30,  
    learning_rate=0.5, random_state=42)  
ada_clf.fit(X_train, y_train)
```

จำนวน predictor
ที่จะสร้าง

Gradient Boosting

- Train predictors sequentially, but later predictors **try to predict residual errors from previous predictors**
- **Learning rate**: how much contribution each predictor gives
 - If low, needs more predictors



In Sklearn

import

```
from sklearn.ensemble import GradientBoostingRegressor  
  
gbrt = GradientBoostingRegressor(max_depth=2, n_estimators=3,  
                                  learning_rate=1.0, random_state=42)  
gbrt.fit(X, y)
```

จำนวน
predictor

Option ของ
decision tree

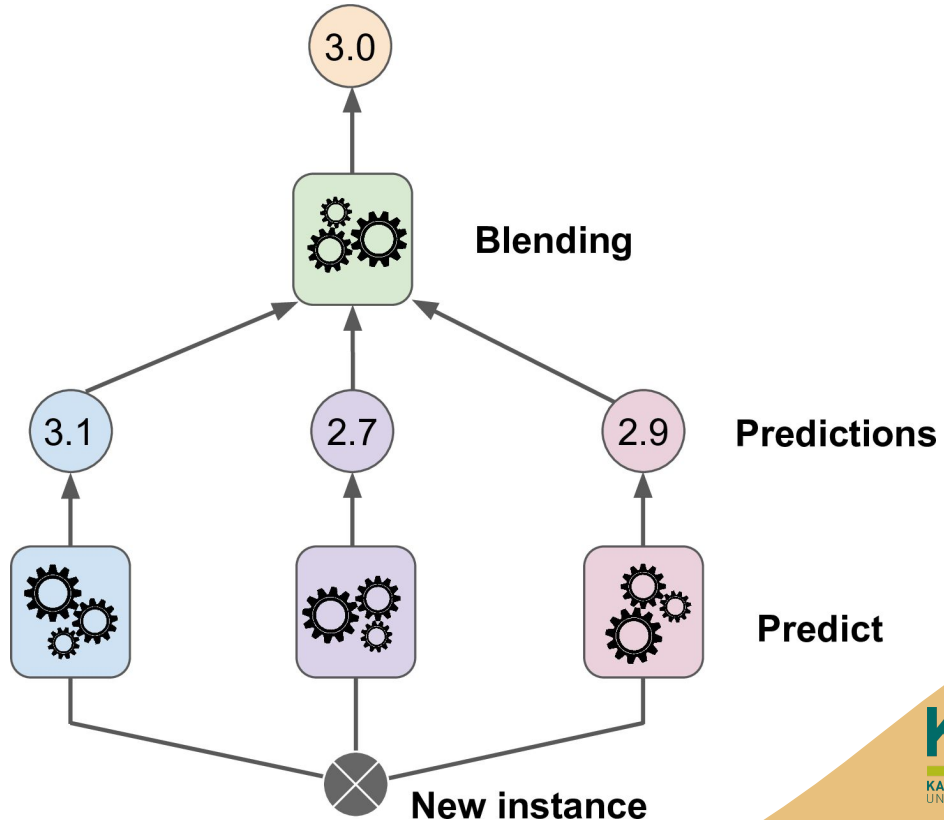
Learning rate

```
gbrt_best = GradientBoostingRegressor(  
    max_depth=2, learning_rate=0.05, n_estimators=500,  
    n_iter_no_change=10, random_state=42)  
gbrt_best.fit(X, y)
```

หยุดการสร้างโมเดลเมื่อการ
ทำนายไม่เปลี่ยนแปลงเท่ากับ
จำนวนรอบที่ระบุ

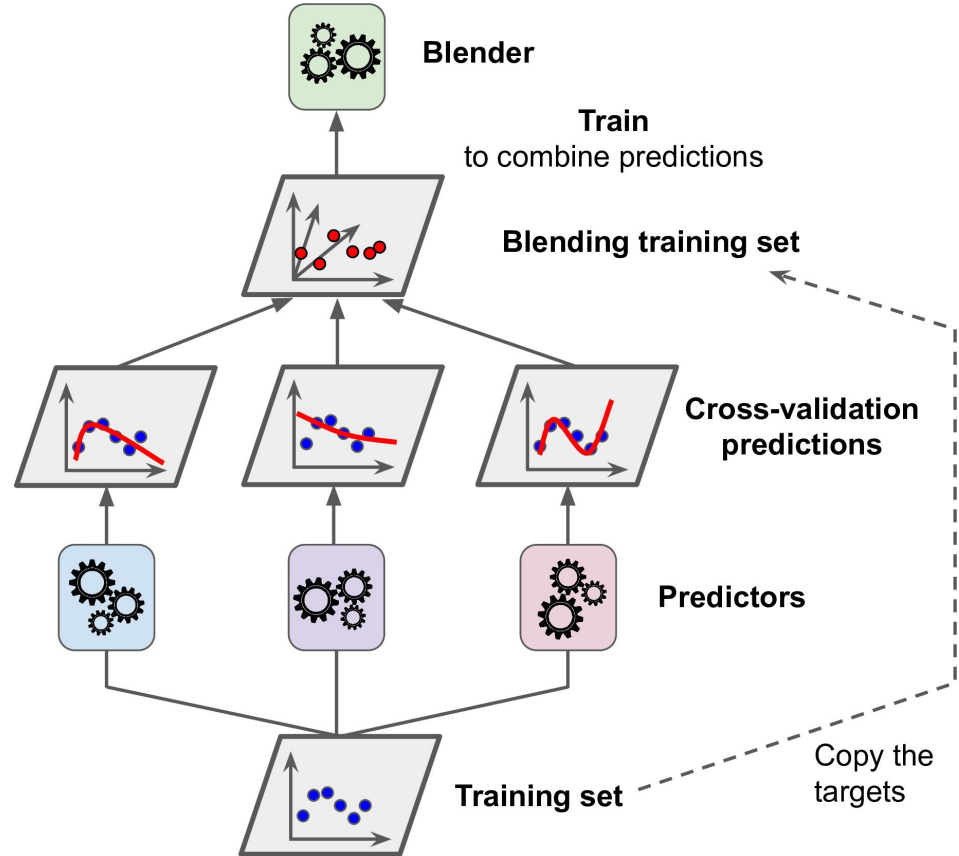
Stacking (stacked generalization)

- Similar to Voting, but train a **blender** or **meta learner** to aggregate predictions of other models



Training a meta learner

- Feature: prediction of each instance made by each predictor (can be obtained using cross-validation)
- Label: the truth value of each instance



In Sklearn

```
from sklearn.ensemble import StackingClassifier

stacking_clf = StackingClassifier(
    estimators=[
        ('lr', LogisticRegression(random_state=42)),
        ('rf', RandomForestClassifier(random_state=42)),
        ('svc', SVC(probability=True, random_state=42))
    ],
    final_estimator=RandomForestClassifier(random_state=43),
    cv=5 # number of cross-validation folds
)
stacking_clf.fit(X_train, y_train)
```

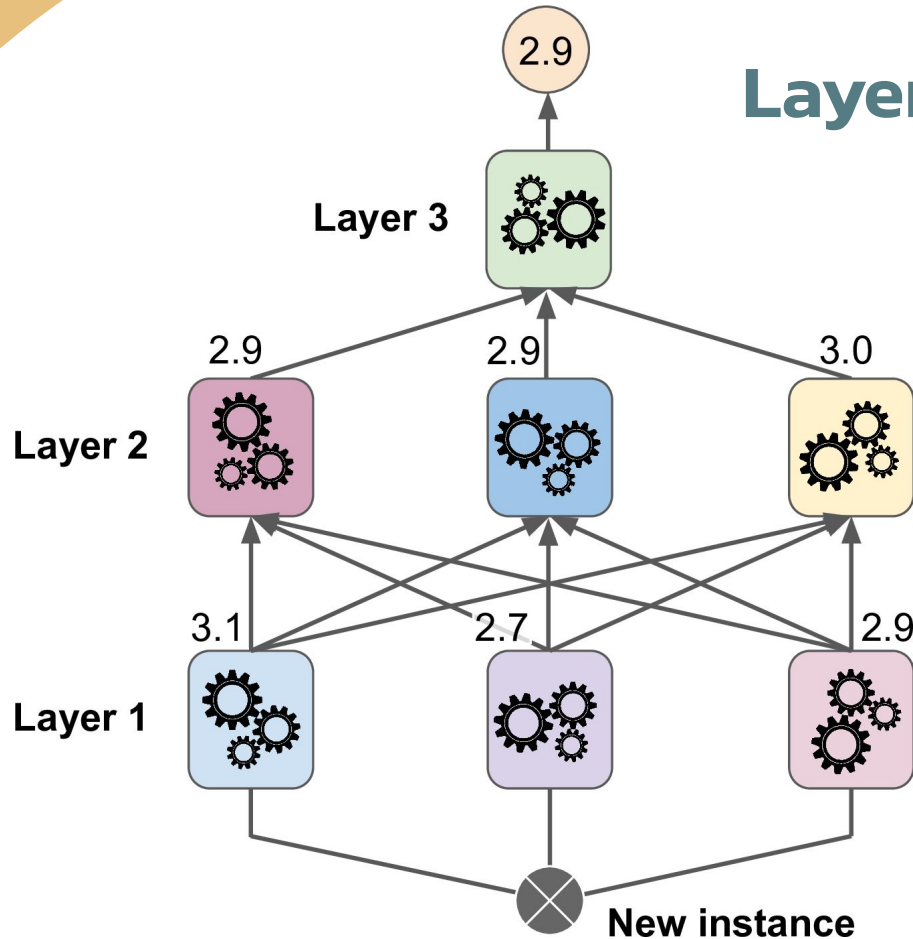
List of
predictors

Each predictor is a
tuple of name and
predictor object

import

Meta learner
object

Layers of Blenders



References

- GALTON, F. "Vox Populi" . Nature 75, 450–451 (1907).
<https://doi.org/10.1038/075450a0>
- Aurélien Géron, "Hands-On Machine Learning with Scikit-Learn and TensorFlow", O'Reilly Media, Inc., March 2017.